

Neuro-Symbolic Probabilistic Programs for Robust 3D Scene Perception

by

Austin J. Garrett

Department of Electrical Engineering and Computer Science
Proposal for Thesis Research in partial fulfillment of the requirements
for the degree of

Master of Engineering in Electrical Engineering and Computer Science
at the

MASSACHUSETTS INSTITUTE OF TECHNOLOGY

December 2019

© Massachusetts Institute of Technology 2019. All rights reserved.

Author
Austin J. Garrett
Department of Electrical Engineering and Computer Science
December 11, 2019

Certified by.....
Vikash K. Mansinghka
Principle Research Scientist, Department of Brain and Cognitive
Sciences
Thesis Supervisor

Neuro-Symbolic Probabilistic Programs for Robust 3D Scene Perception

by

Austin J. Garrett

Submitted to the Department of Electrical Engineering and Computer Science
on December 11, 2019, in partial fulfillment of the
requirements for the degree of
Master of Engineering in Electrical Engineering and Computer Science

Abstract

In this thesis, I consider Bayesian approaches to doing inference over structured scenes. This includes the design and implementation of a probabilistic language for representing semantic relations between multiple objects. I propose a strategy for integrating neural techniques into MCMC-based techniques for inverse graphics, where a generative model is constructed to predict the behavior of neural detectors, allowing for semantic semantic structuring of the output. I further discuss possible experimentation with a mixed likelihood model that combines the strengths of top-down scene structuring with analysis-by-synthesis approaches for inverse graphics. I finally discuss the development of a framework for realtime particle filter based inference based on the Robot Operating System and Gen probabilistic programming language, as well as its integration into the broader Cora project for probabilistic modeling of intuitive physics.

Thesis Supervisor: Vikash K. Mansinghka

Title: Principle Research Scientist, Department of Brain and Cognitive Sciences

Chapter 1

Thesis Proposal

1.1 Introduction

The past decade has seen a flourishing of highly powerful computer vision techniques, due to the success of deep learning techniques accelerated on massively parallel hardware.

Yet the shortcomings of neural techniques have only grown increasingly urgent in proportion to their dominance. While these approaches have resulted in the amortization of traditionally very difficult and high-dimensional computer vision problems, the lack of semantic structuring in their output has led to great challenges in their deployment in practical domains where interpretability and composability are desperately needed to integrate these components into larger frameworks.

Meanwhile probabilistic techniques offer a more principled approach to modeling structured data, and the ecosystem for using these tools has rapidly expanded. Such environments are desirable for their expressibility, corresponding to the fact that they often represent intuitive models of rationality inspired from cognitive and neuroscience which increasingly rely on such modeling to explain the human reasoning. In this thesis, I will explore probabilistic methods for combining neural detectors with richly structured priors in order to filter and improve the robustness of the network's output.

1.2 Prior and Related Works

1.2.1 Accelerating Inverse Graphics

An alternative perspective on the problem of pose estimation comes from the perspective of “analysis-by-synthesis”, in which a generative model allows for MCMC-based inference, integrating a notion of uncertainty and robustness into the estimation procedure. Typically to make inference efficient, these approaches include some amount of discriminative training to amortize the inference procedure with data-driven kernels. *Picture* is a probabilistic programming language for scene perception that proposed among others, arbitrary blocked Gibbs moves, gradient-based proposals, and elliptical slice sampling in order to incorporate bottom-up amortization of MCMC [2]. Jampani et. al. created the *informed sampler*, a mixed inference kernel between a trained state-independent discriminator proposal, and a local Gaussian metropolis-hastings move [1]. There has also been work on directly using neural networks as initialization for MCMC. Yildirim et. al. leveraged convolutional neural networks to initialize a Markov chain for face processing [6].

1.2.2 Neural Approaches to Pose Estimation

Only recently have neural approaches successfully begun to make headway on full 6D (3 spatial, 3 rotational) pose estimation for multiple objects. These algorithms offer a powerful modeling perspective that can be used as an intermediate step for segmentation, bounding box regression, or feature extraction. However, they often suffer from the same brittleness and lack of insight as deep learning in general.

[5] combines heuristical Hough voting and convolutional feature extraction to approach the problem. [4] uses a neural network to estimate belief maps of keypoints in 2D image coordinates that are fed into a standard perspective- n -point (PnP) algorithm to recover the full 6D pose. It was trained only on synthetic data, using domain randomization to generalize to the real world. [3] uses a render-and-compare based loss inspired by inverse graphics methods to train the neural regressor.




Figure 1-1: Two identical representations of three blocks stacked on top of a floor in a virtual environment. The right side assumes no semantic relational information, with 24 parameters required to specify the positions of all the objects in the scene. The left side represents semantic information about the relationship between objects, namely edges represent an “on-top of” relationship. Under this scene structure, only 15 parameters are necessary for specifying the same scene.[7]

1.2.3 Bayesian Inference over Structured Scenes

Extending from this previous work on inverse graphics, we wish to construct generative models that contain richer *semantic information* about objects in the scene – namely human-interpretable information about the relations between objects. This can be crucial in reducing the complexity of inference by reducing the dimensionality of the parametrization to be inferred.

Scene Graph Representation

To encode structural information, we leverage a common computer vision representation called a *scene graph*. In this graph, poses are represented as transformations from object to object, or from implicit world frame to object. Directed edges specify the type of relative pose, which encodes information about the way objects are geometrically situated. For example, an edge may represent “contact”, which requires specifying a contact plane, a 2D translation, and a single rotational DoF.

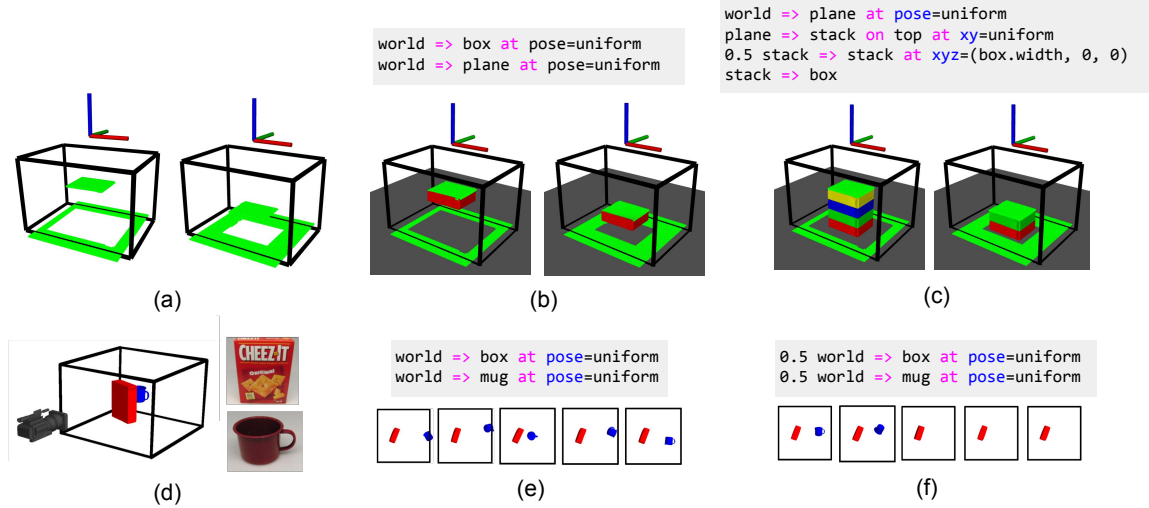


Figure 1-2: Two scenarios with prior knowledge specified as programs in our probabilistic scene description language. (a) shows two depth measurements made by a depth camera. (b) shows a program in which boxes have random poses in 3D space, and the resulting inferred pose of the box. (c) shows a program that assumes boxes are in stacks that rest on the floor, and the inferred number of boxes that explain the data for each observation. (d) shows another scenario, where a depth camera observes a box that is occluding a mug. (e) shows a program that asserts that both objects exist (but at unknown poses). The resulting inferences show the mug must exist somewhere behind the box to be consistent with this knowledge and the observation. (f) shows a program that allows for either object to not exist, and the resulting joint inferences about the mug’s existence and pose.

Probabilistic Scene Description Languages

It can often be desirable to specify knowledge in a programmatic way. One way to define a generative model over a scene graph representation is through a probabilistic context free grammar called a *scene description language*. Such a language can further increase the expressibility of the scene graph concept by adding existential uncertainty to the objects themselves. This is encoded by some probability over an expansion of the production rules in the PCFG.

Even without the use of data-driven proposals, we can demonstrate how this language for uncertain knowledge representation can recover common-sense reasoning that naturally combines prior knowledge with observation to obtain posteriors with rich structural information.




Figure 1-3: Neural detections can be inaccurate, and violate semantic relations between objects (eg. allow for object interpenetration). Given a neural detector, we can treat the flat and noisy detections as observations under a generative model, given prior structural information about the scene. Thus we implicitly model and correct for the failure modes of neural detections using uncertain prior structural knowledge. Using custom inference kernels, we can potentially recover the network’s uncertainty, resolve impossible scenarios that violate prior knowledge, or even recover qualitative pose relationships like contact (green arrow) to enrich neural detections.

1.3 Proposed Work

| Month | Work to be completed |
|----------|--|
| December | Thesis proposal and ROS-based infrastructure for online particle filtering and associated visualization. |
| January | Experimentation with observational models for modeling neural detections. Writing submission to the RSS conference. |
| February | Engineering detectors to be trained on synthetic data. |
| March | Research on robust mixture likelihood; including the rendering-based generative model to solidify “meta-inference” perspective, where the neural detector is a <i>modeled</i> amortization of the inference procedure. |
| April | NeurIPS experiments (model-based inference cleanup/error correction, training detectors using generative model). |
| May | Finalizing NeurIPS experimentation; writing submission. |
| June | Translation of contact-graph representation into Cora project. |
| July | Refactoring code base and integration of the ROS particle filter and visualization components into Cora. |
| August | Finalization of thesis and final documentation of work and code base. |

1.3.1 Neuro-predictive Generative Modeling

Modeling a rendering pipeline in a fully Bayesian way suffers from certain computational challenges. While guaranteed to converge to a correct posterior with unlimited compute time, the non-asymptotic properties of MCMC are poorly-understood. Our preliminary work in inverse graphics techniques applied to structured scenes sup-

ports the hypothesis that naive analysis-by-synthesis approaches that only leverage a generative rendering pipeline often fail to explore all the modes of the posterior in reasonable time bounds.

Concretely, for scene graph \mathcal{S} sampled from from a structured prior program like those discussed in 1.2.3, and continuous parameterization $\vec{\nu}_{\mathcal{S}}$, a rendering-based likelihood relies on modeling the image data using a full rendering pipeline $R(\mathcal{S}, \vec{\nu}_{\mathcal{S}}) = X$, where X is a rendered image. For robustness, the likelihood is modeled not directly on X , but on a noisy function of the pixel data. The noise is modeled is a mixture of a uniform distribution on the range of possible depth values $[0, D]$ and a normal distribution with mean centered at the true rendered pixel value $R(\mathcal{S}, \vec{\nu}_{\mathcal{S}})$ with fixed variance σ^2 , leaving the full likelihood on noisy image data Y as

$$p(Y|\mathcal{S}, \vec{\nu}_{\mathcal{S}}) = \prod_{r=1}^R \prod_{c=1}^C \left(0.1 \cdot \frac{1}{D} + 0.9 \cdot \mathcal{N}(Y_{r,c}; R(\mathcal{S}, \vec{\nu}_{\mathcal{S}})_{r,c}, \sigma) \right) \quad (1.1)$$

1.1 is very high-dimensional, making it prone to capture by local minima. Neural networks empirically demonstrate strong performance at locating strong *maximum a posteriori* estimates in bounded compute time. Thus, we may consider combining neural techniques with MCMC to improve performance of inference. One perspective may consider neural networks as amortizations of the inference procedure. Consider a neural detector $\phi_{\mathcal{S}}$ that outputs poses under a given parametrization \mathcal{S} (often a full unconstrained 6D pose), such that ideally $\phi_{\mathcal{S}}(R(\mathcal{S}, \vec{\nu}_{\mathcal{S}})) = \vec{\nu}_{\mathcal{S}}$. Given image data Y , we can use the corresponding estimate $\phi_{\mathcal{S}}(Y)$ as an initialization for some MCMC technique.

However, this method is rather unprincipled in that it uses a point estimate without uncertainty quantification, thus implicitly relying on $\phi_{\mathcal{S}}$ robustly predicting the mode of the distribution. This assumption is not guaranteed, and in fact when neural networks do fail, they often fail catastrophically, estimating wildly incorrect poses. Furthermore, because it is only a heuristic and not a full proposal distribution, it is not obvious how to combine neural detections across multiple timesteps into a coherent picture. The issue fundamentally lies in the fact that this approach contains no

way to reason about the behavior of the bottom-up proposals. If we wish to robustly use these models, we need to leverage information about their behavior.

Crucially, we can observe that the failure modes of these neural techniques are often quite predictable with the use of certain easy to compute statistics. When objects are too close or too far away, at strange angles, or heavily occluded, the neural detector is much more likely to fail. To account for these configurations, we propose instead modeling the neural detections as observations to a generative model that attempts to retrieve the latent scene graph \mathcal{S} that predicts the detections from the neural network $\phi_{\mathcal{S}}(Y)$.

A minimal observational model may be a simple mixture between a Gaussian and a uniform distribution over the volume V of the whole space

$$p(\phi_{\mathcal{S}}(Y)|\mathcal{S}, \vec{\nu}_{\mathcal{S}}) = 0.1 \cdot \frac{1}{V} + 0.9 \cdot \mathcal{N}(\phi_{\mathcal{S}}(Y); T_{6D}(\vec{\nu}_{\mathcal{S}}), \sigma) \quad (1.2)$$

Where $T_{6D}(\vec{\nu}_{\mathcal{S}})$ is a projection of $\vec{\nu}_{\mathcal{S}}$ into a 6D pose relative to the camera. Even such a simple observational model can be sufficient to filter some noise from the bottom-up proposal, by roughly modeling that the neural network sometimes makes mistakes (although in this case we don’t use information about whether this is more or less likely at any given time).

1.3.2 Neural Predictor Experimentation

The first task concrete task is creating the infrastructure sufficient for creating a minimal demo. For this we will develop an end-to-end system that uses DOPE+Gen+ROS to track a fixed set of objects, integrating prior knowledge about temporal consistency of object trajectories via state space modeling. For this, we can use a simple particle filter over object trajectories. This will help us nail down the end-to-end integration aspect, and also serves as the minimal example of “filtering the output of a neural network via prior knowledge”. For simplicity we will initially work with an observation model of the form 1.2. Even with this simple model, we hypothesis it is possible to filter spurious neural detections using physical assumptions of object persistence and

inertial trajectories. A very simple experiment along these lines may consist of an object passing behind an occluder. As the object becomes more and more occluded, the quality of the neural detections corresponding to the object become poorer. A successful experiment would demonstrate our technique is able to smooth out spurious detections, and maintain beliefs of the object’s existence once it is no longer visible in the scene.

Additional experimentation would use the full pipeline to infer scene structure, integrating prior knowledge about scene graphs (e.g. contact relationships) as well as temporal consistency. In this version we may construct an experiment where an object rests on another object, and test if we can successfully recover the “contacting” relationship between objects. If this is successful, we may then consider pushing the supporting object around and test if we can maintain stability in the relationship type over time.

1.3.3 Robust Mixture Likelihood Experimentation

A natural question extending from neural modeling is how it relates back to the rendering-based generative approach we initially considered. While simply modeling the inverse procedure can be sufficient for decreasing confidence in inaccurate detections, to increase the actual accuracy it is very useful to step the neural detections with respect to how well it explains the bottom-level data. Thus we might consider an extension where we reintroduce the rendering generative model from intermediate flat detections to the actual observation.

In this case instead of modeling just the neural network component, we also can leverage information from the inference procedure. This can give us insight into whether or not the neural initialization is stable with respect to the image data; if we find high variance in MCMC moves we can use this information to infer the detections were poor explanations for the data, whereas if the moves are tightly bound in a localized region we may increase our confidence in the quality of the neural detections.

This experimentation may contrast with the work expected in January; we may

demonstrate that given a rendering based-likelihood we can not only decrease confidence in inaccurate predictions, but we can increase the overall accuracy of the neural detections. For example, we may find a scene in which neural detections correctly estimate position, but fail to localize rotation appropriately (a common failure mode of neural detections), and demonstrate that the robust mixture likelihood can recover the true rotation by first recovering a possible qualitative relationship (eg. object is on-top of a table), and then localize rotation on the reduced parametrization.

Additional experiments may include training the neural detector using samples from the scene graph generator. In this way we can enrich the training data with synthetic scenes, and increase the robustness of the detector when operating on samples that are well-captured by the structural prior.

1.3.4 Engineering Infrastructure

All of this experimentation requires sophisticated and novel infrastructure for the prototyping of complex custom inference kernels and observational models to be applied in realtime. The domain of realtime scene perception is heavily related to robotics applications, and thus the Robot Operating System serves as a helpful abstraction for the various components that go into our infrastructure. Namely, we assume asynchronous operation of multiple components of the overall infrastructure, corresponding to sensors, the inference backend, and visualization. These components will contain several reusable artifacts that will be integrated into the broader Cora project near the end of the Master's.

Scene Graph Representation

A crucial component of the Cora project is the scene graph representation. The representation used for our experimentation will be in concurrent development with that of the Cora project, ideally allowing us to rapidly refactor and transfer our inference techniques into the Cora project.

Realtime Particle Filtering

The inference node receives observations from sensors as published messages from sensors in realtime. We will support both online inference (dropping observations if they come during particle filter updates), and offline inference (queueing observations and sequentially processing them with as much time as necessary). The former approach supports the online demands of robotics applications, while the latter can be used for unconstrained experimentation with the inference algorithm.

Visualization

Visualization of the inference procedure is crucial for gaining intuition of the failure modes and rates of convergence. In particular, it is essential to be able to view the particle filter beliefs in realtime. In our framework, we will use a visualization based on ROS' asynchronous RViz visualizer. The inference node can actively publish its beliefs to this node, creating a powerful interface for peering into the internals of various inference kernels. The Cora project has large overlaps in demands for inference, and so this tool will be a highly useful artifact for the project overall.

1.4 Conclusion

In this thesis, I will consider methods for integrating uncertain symbolic information with bottom-up neural proposals. We have already experimented with a robust rendering-based likelihood model without neural proposals, demonstrating some limited results on recovering common-sense scene understanding. My next steps include using this symbolic information to filter neural detections by constructing a generative model to represent the possibly spurious behavior of the detector. After this I will consider methods for reintegrating the rendering-based likelihood into a robust mixture model in order to increase the accuracy of neural detections using both structural information and analysis-by-synthesis approaches. Finally, I will integrate essential components of the infrastructure into the broader Cora project so they may be reused by other researchers.

Bibliography

- [1] Varun Jampani, Sebastian Nowozin, Matthew Loper, and Peter V. Gehler. The informed sampler: A discriminative approach to bayesian inference in generative computer vision models. *CoRR*, abs/1402.0859, 2014.
- [2] Tejas D Kulkarni, Pushmeet Kohli, Joshua B Tenenbaum, and Vikash Mansinghka. Picture: A probabilistic programming language for scene perception. In *Proceedings of the ieee conference on computer vision and pattern recognition*, pages 4390–4399, 2015.
- [3] Abhijit Kundu, Yin Li, and James M Rehg. 3d-rcnn: Instance-level 3d object reconstruction via render-and-compare. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, pages 3559–3568, 2018.
- [4] Jonathan Tremblay, Thang To, Balakumar Sundaralingam, Yu Xiang, Dieter Fox, and Stan Birchfield. Deep object pose estimation for semantic robotic grasping of household objects. *CoRR*, abs/1809.10790, 2018.
- [5] Yu Xiang, Tanner Schmidt, Venkatraman Narayanan, and Dieter Fox. Posecnn: A convolutional neural network for 6d object pose estimation in cluttered scenes. *arXiv preprint arXiv:1711.00199*, 2017.
- [6] Ilker Yildirim, Tejas D Kulkarni, Winrich A Freiwald, and Joshua B Tenenbaum. Efficient analysis-by-synthesis in vision: A computational framework, behavioral tests, and comparison with neural representations. In *Proceedings of the thirty-seventh annual conference of the cognitive science society*, 2015.
- [7] Ben Zinberg, Marco Cusumano-Towner, and Vikash K. Mansinghka. Structured differentiable models of 3d scenes via generative scene graphs. *Workshop on Perception as Generative Reasoning, NeurIPS 2019, Vancouver, Canada.*, 2019.